

Diving Deep into Bedrock AgentCore

Diving Deep into Bedrock AgentCore

► Prerequisites

Amazon Bedrock AgentCore Fundamentals

▼ AgentCore Runtime

Hosting Runtime Agent

Hosting MCP server

Advanced Concepts for Agentcore Runtime

► AgentCore Gateway

► AgentCore Identity

▼ AgentCore Memory

AgentCore Memory: Short-Term Memory

AgentCore Memory: Long-Term Memory Strategies

▼ AgentCore Tools

AgentCore 1P Tool - AgentCore Code Interpreter

▼ AgentCore 1P Tool - AgentCore Browser

AgentCore Browser - Nova Act SDK

AgentCore Browser - Browser-Use

▼ AgentCore Observability

AgentCore Observability for Runtime hosted Agent

Observability for Non-Runtime hosted Agents

AgentCore ↗

AgentCore Documentation ↗

► AWS account access

Workshop catalog in AWS Builder Center ↗

► Content preferences

Exit event

Event ends in 17 hours 59 minutes.

[Event dashboard](#) > [AgentCore Runtime](#) > [Hosting Runtime Agent](#)

Hosting Runtime Agent

Overview

[Amazon Bedrock AgentCore Runtime ↗](#) enables developers to deploy and run AI Agents in a secure and scalable serverless environment. The Bedrock AgentCore Runtime provides a secure, serverless hosting environment necessary for deploying and running AI agents and tools without managing infrastructure. By offering a framework-agnostic and model-flexible runtime, the AgentCore Runtime enables organizations to leverage their existing agent code and models, rather than being limited to a specific development framework or AI architecture.

How it works

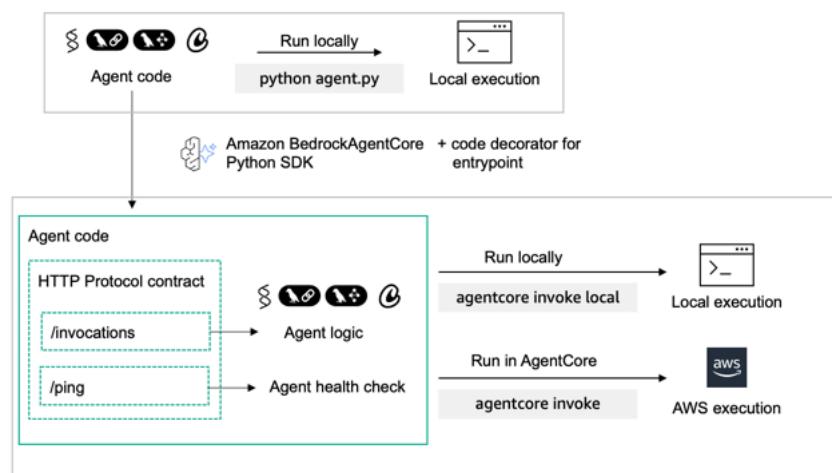
The Amazon Bedrock AgentCore Python SDK acts as a wrapper that:

- **Transforms** your agent code into AgentCore's standardized protocols
- **Handles** HTTP and MCP server infrastructure automatically
- **Lets you focus** on your agent's core functionality
- **Supports** two protocol types:
 - **HTTP Protocol**: Traditional request/response REST API endpoints
 - **MCP Protocol**: Model Context Protocol for tools and agent servers

High Level Architecture

When hosting agents, the SDK automatically:

- Hosts your agent on port 8080
- Provides two key endpoints:
 - /invocations: Primary agent interaction (JSON input → JSON/SSE output)
 - /ping: Health check for monitoring



Getting Started

First, you need to implement an AI Agent using the framework of your choice. In this example, we implement a simple AI Agent using the [Strands Agent](#)  framework:

```

1   from strands import Agent, tool
2   from strands_tools import calculator
3   import argparse
4   import json
5   from strands.models import BedrockModel
6
7   @tool
8   def weather():
9       """ Get the weather """
10      return "sunny"
11
12      model_id = "us.anthropic.claude-3-7-sonnet-20250219-v1:0"
13      model = BedrockModel(
14          model_id=model_id,
15      )
16      agent = Agent(
17          model=model,
18          tools=[calculator, weather],
19          system_prompt="You're a helpful assistant. You can perform simple math calculations and tell
20      )
21
22  def strands_agent_bedrock(payload):
23      """
24      Invoke the agent with a payload
25      """
26      user_input = payload.get("prompt")
27      response = agent(user_input)
28      return response.message['content'][0]['text']
29
30  if __name__ == "__main__":
31      parser = argparse.ArgumentParser()
32      parser.add_argument("payload", type=str)
33      args = parser.parse_args()
34      response = strands_agent_bedrock(json.loads(args.payload))
35      print(response)

```



Prepare the Agent for AgentCore Runtime

The [Amazon Bedrock AgentCore Python SDK](#)  provides a lightweight wrapper that helps you deploy your agent functions as HTTP services compatible with Amazon Bedrock AgentCore.

You can convert **your existing agent function** into an Amazon Bedrock AgentCore-compatible service with just **four steps**:

1. Import the Runtime App with `from bedrock_agentcore.runtime import BedrockAgentCoreApplication`
2. Initialize the App in our code with `app = BedrockAgentCoreApplication()`
3. Decorate the invocation function with the `@app.entrypoint` decorator
4. Let AgentCoreRuntime control the running of the agent with `app.run()`

```
# strands_claude.py
from strands import Agent, tool
from strands_tools import calculator # Import the calculator tool
import argparse
import json
from bedrock_agentcore.runtime import BedrockAgentCoreApp ←
from strands.models import BedrockModel

app = BedrockAgentCoreApp() ←

# Create a custom tool
@tool
def weather():
    """ Get weather """ # Dummy implementation
    return "sunny"

model_id = "us.anthropic.claude-sonnet-4-20250514-v1:0"
model = BedrockModel(
    model_id=model_id,
)
agent = Agent(
    model=model,
    tools=[calculator, weather],
    system_prompt="You're a helpful assistant. You can do simple math calculation, and tell the weather."
)

@app.entrypoint
def strands_agent_bedrock(payload):
    """
    Invoke the agent with a payload
    """
    user_input = payload.get("prompt")
    print(f"User input: {user_input}")
    response = agent(user_input)
    return response.message['content'][0]['text']

if __name__ == "__main__":
    app.run() ←
```

Invoke your Agent

You can invoke the agent using the `InvokeAgentRuntime` operation:



```
1  import boto3
2  import json
3
4  # Initialize the Bedrock AgentCore client
5  agent_core_client = boto3.client('bedrock-agentcore', region_name="us-east-1")
6
7
8  # Prepare the payload
9  payload = json.dumps({"prompt": prompt}).encode()
10
11 # Invoke the agent
12 response = agent_core_client.invoke_agent_runtime(
13     agentRuntimeArn=agent_arn,
14     runtimeSessionId=session_id,
15     payload=payload
16 )
```

Try it out

At an AWS Event

In the JupyterLab UI, navigate to `01-AgentCore-runtime/01-hosting-agent/01-strands-with-bedrock-model`

Self-paced

Here's are notebooks that lets you try out the above and extend the patterns to other frameworks and models

Example	Framework	Model	Description
strands-with-bedrock-model ↗	Strands Agents	Amazon Bedrock	Basic agent hosting with AWS native models
langgraph-with-bedrock-model ↗	LangGraph	Amazon Bedrock	LangGraph agent workflows

Example	Framework	Model	Description
strands-with-openai-model ↗	Strands Agents	OpenAI	Integration with external LLM providers

[Previous](#)[Next](#)